

# Feature Engineering

Dmitry Larko, Sr. Data Scientist @ H2O.ai

dmitry@h2o.ai

# About me



## Dmitry Larko

Sr. Data Scientist at H2O.ai

San Francisco Bay Area, CA, United States

Joined 5 years ago · last seen in the past day

[Twitter](#) [LinkedIn](#) <https://h2o.ai>

Followers 59



Competitions  
Grandmaster

[Home](#) [Competitions \(37\)](#) [Kernels \(0\)](#) [Discussion \(32\)](#) [Datasets \(0\)](#) ...

[Edit Profile](#)

### Competitions Grandmaster



Current Rank

69

of 66,441

Highest Rank

25



9



11



6

[Grupo Bimbo Inventory De...](#)

🥇 · a year ago · Top 1%

1<sup>st</sup>

of 1969

[Walmart Recruiting II: Sale...](#)

🥇 · 3 years ago · Top 1%

2<sup>nd</sup>

of 485

[Acquire Valued Shoppers C...](#)

🥇 · 3 years ago · Top 1%

2<sup>nd</sup>

of 952

### Kernels Contributor



Unranked



0



0



0

No kernel results

### Discussion Contributor



Unranked



0



4



12

[How to Tuning XGboost in ...](#)

🥇 · 2 years ago

17

votes

[How to Tuning XGboost in ...](#)

🥇 · 2 years ago

14

votes

[Data Scientist Hero](#)

🥇 · 2 years ago

8

votes

H<sub>2</sub>O.ai

# Feature Engineering

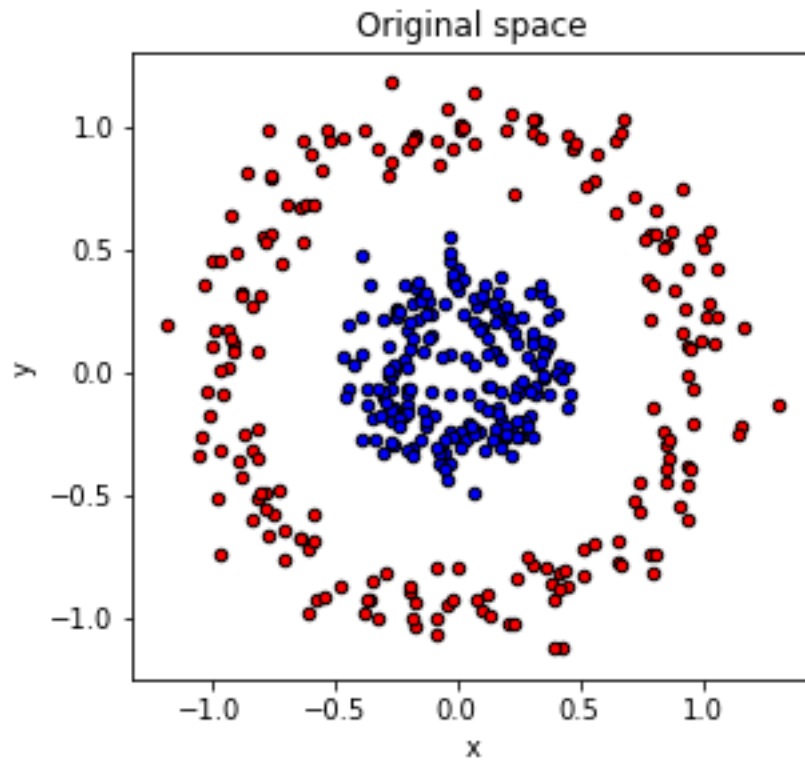
"Coming up with features is difficult, time-consuming, requires expert knowledge. "Applied machine learning" is basically feature engineering." ~ Andrew Ng

"... some machine learning projects succeed and some fail. What makes the difference? Easily the most important factor is the features used." ~ Pedro Domingos

"Good data preparation and feature engineering is integral to better prediction" ~ Marios Michailidis (KazAnova), Kaggle GrandMaster, Kaggle #3, former #1


"*you have to turn your inputs into things the algorithm can understand*" ~ Shayne Miel, answer to ["What is the intuitive explanation of feature engineering in machine learning?"](#)

# What is feature engineering




Not possible to separate using linear classifier

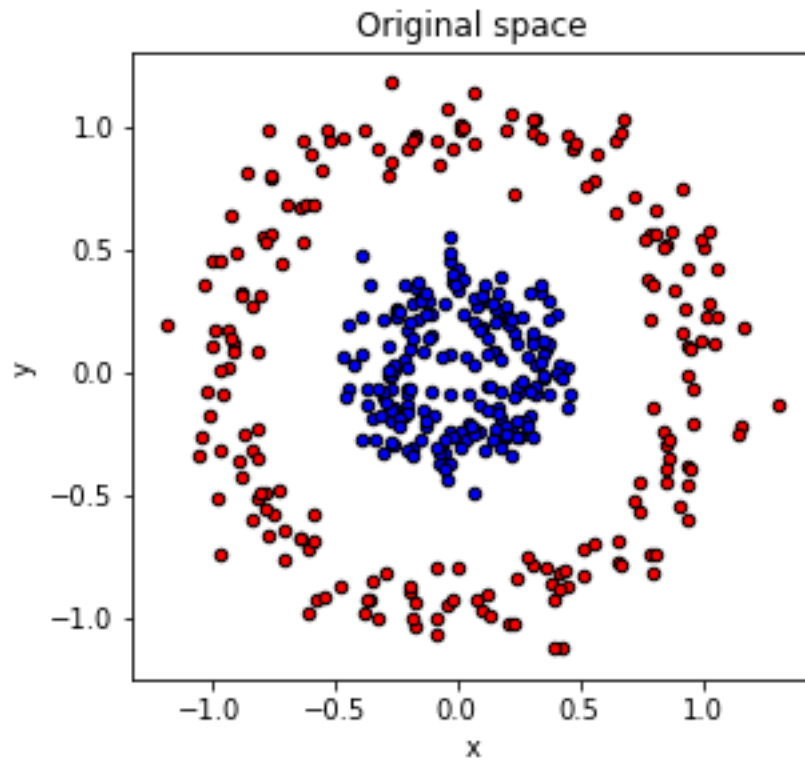
# What is feature engineering




What if we use polar coordinates instead?



# What is feature engineering




Cartesian coordinates




$$P = (X, Y)$$

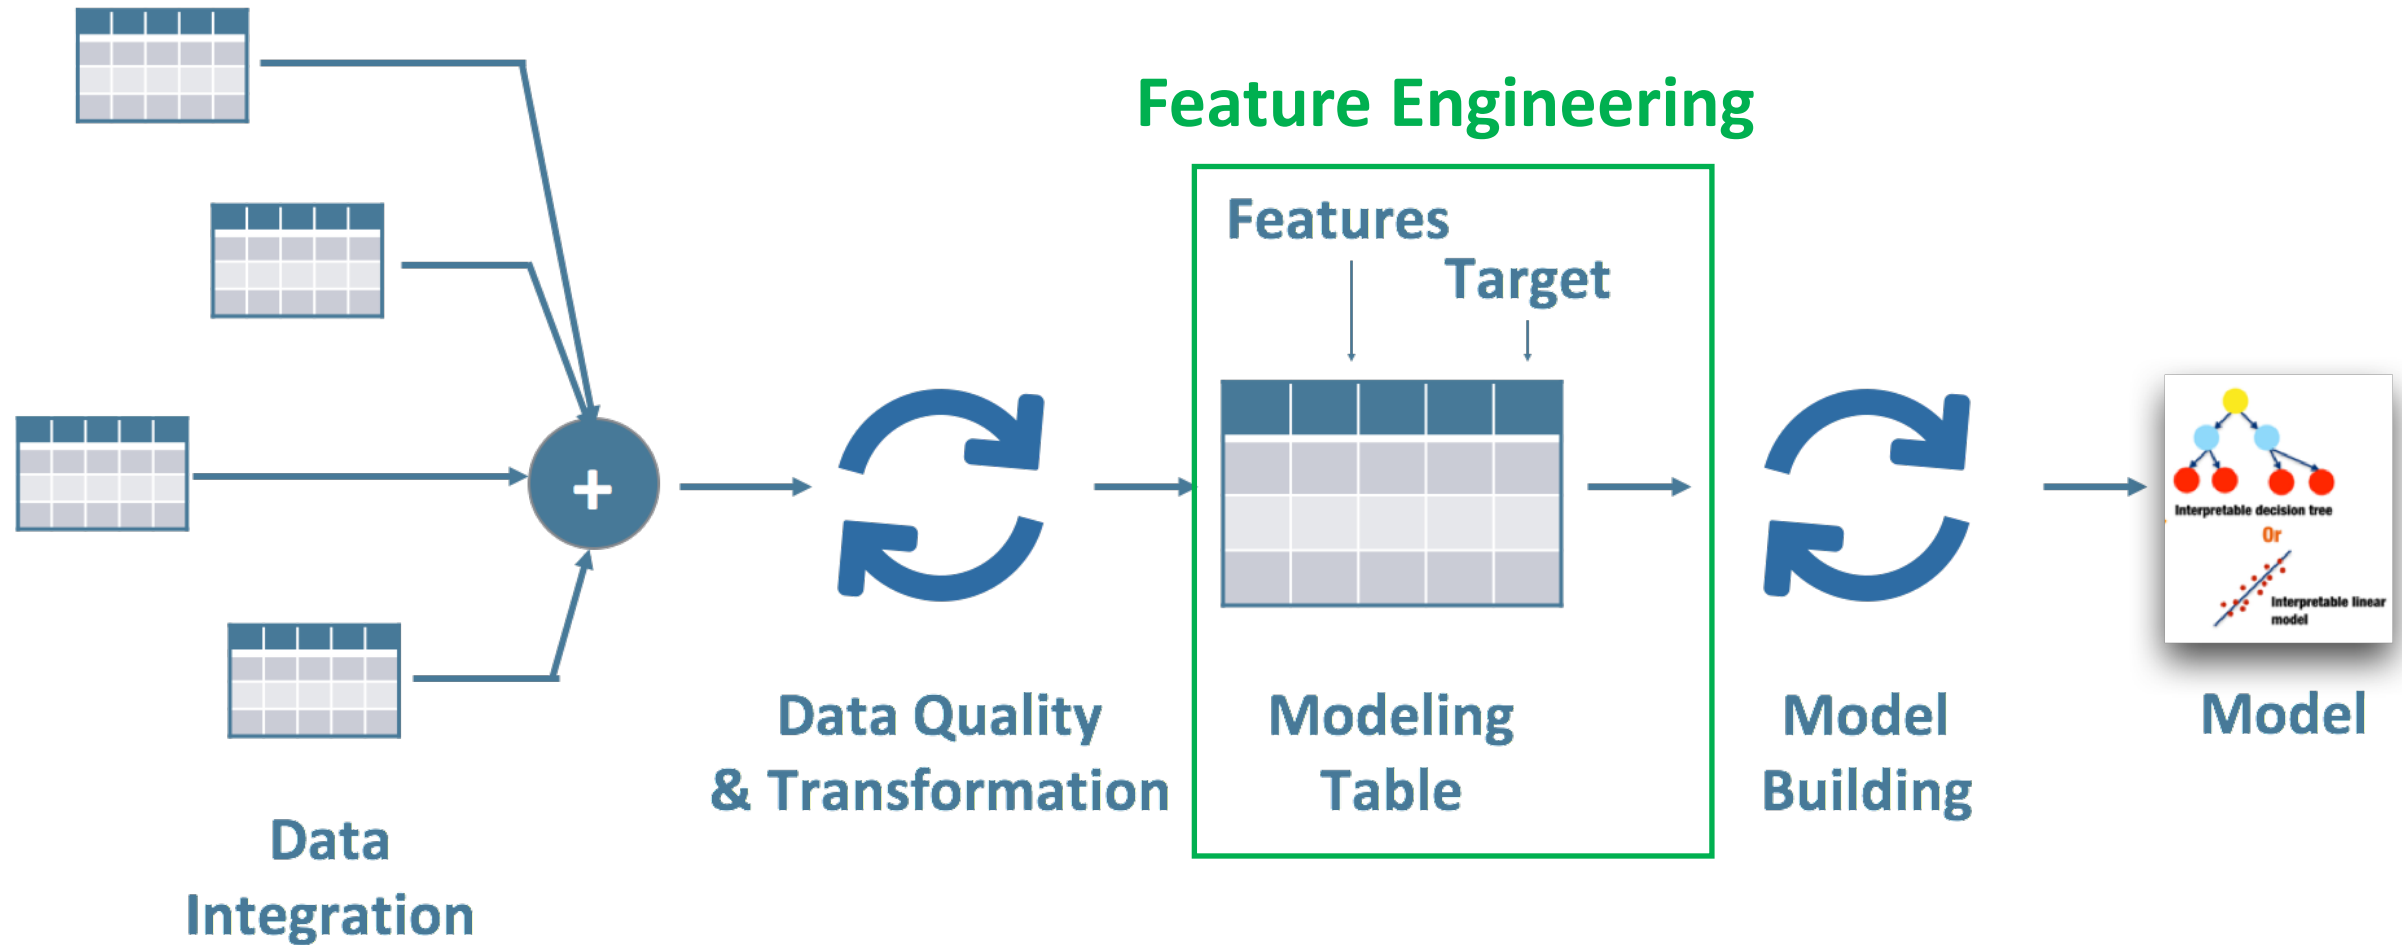
Polar coordinates



$$P = (\theta, r)$$



# Typical Enterprise Machine Learning Workflow

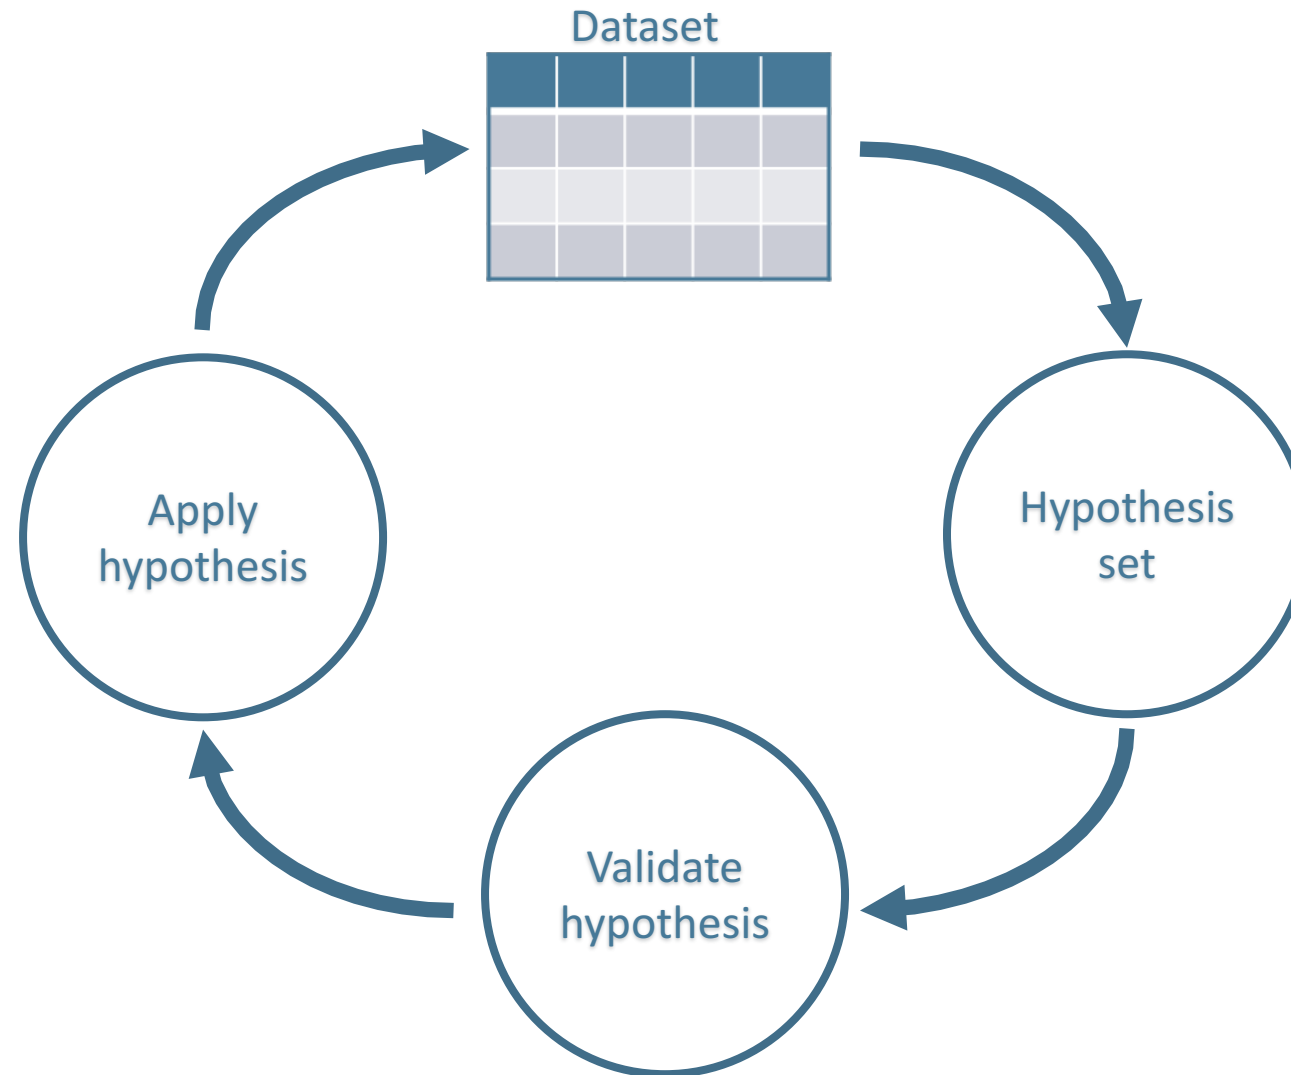


# That's NOT Feature Engineering

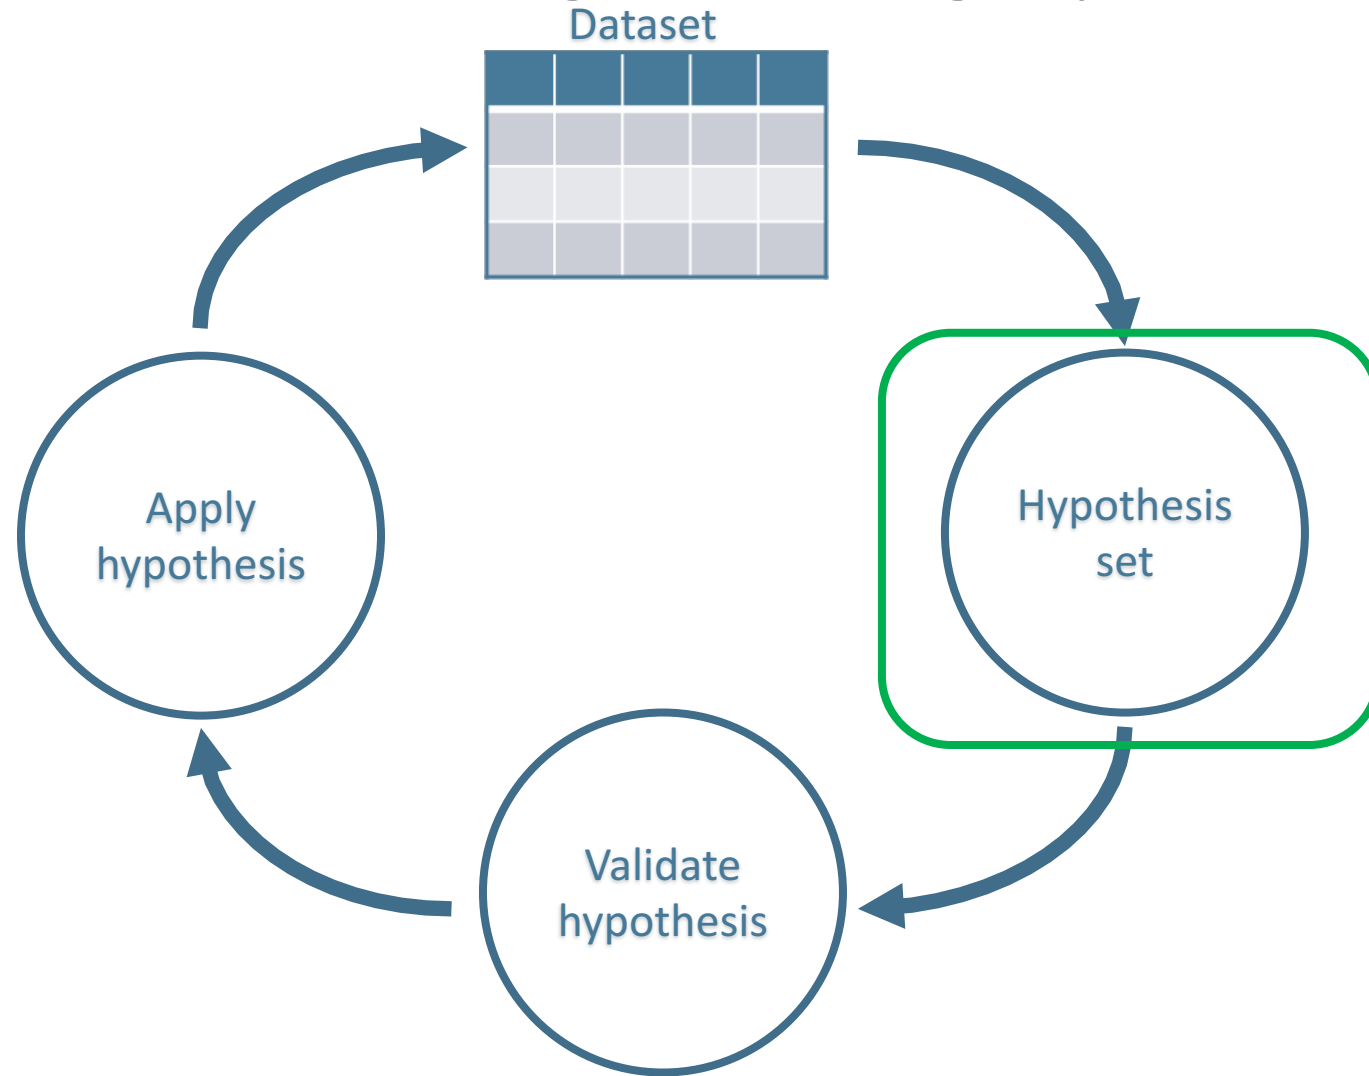
- Initial data collection
- Creating the target variable.
- Removing duplicates, handling missing values, or fixing mislabeled classes, it's data cleaning.
- Scaling or normalization
- Feature selection, although I'm going to mention it in this talk 😊



# Feature Engineering cycle




# Feature Engineering cycle



How?

- Domain knowledge
- Prior experience
- EDA
- ML model feedback

# Feature Engineering cycle



How?

- Cross-validation
- Measurement of desired metrics
- Avoid leakage

# Why feature engineering is hard

- Powerful feature transformations (like target encoding) can introduce leakage when applied wrong
- Usually requires domain knowledge about how features interact with each other
- Time-consuming, need to run thousand of experiments

# Feature Engineering


- Extract more new gold features, remove irrelevant or noisy features
  - Simpler models with better results
- Key Elements
  - Target Transformation
  - Feature Encoding
  - Feature Extraction

# Target Transformation

- Predictor/Response Variable Transformation
  - Use it when variable shows a skewed distribution make the residuals more close to “normal distribution” (bell curve)
  - Can improve model fit
    - $\log(x)$ ,  $\log(x+1)$ ,  $\sqrt{x}$ ,  $\sqrt{x+1}$ , etc.

# Target Transformation

In Liberty Mutual Group: Property Inspection Prediction



Different transformations might lead to different results

# Feature Encoding

- Turn categorical features into numeric features to provide more fine-grained information
  - Help explicitly capture non-linear relationships and interactions between the values of features
  - Most of machine learning tools only accept numbers as their input
    - xgboost, gbm, glmnet, libsvm, liblinear, etc.



# Feature Encoding

- Labeled Encoding
  - Interpret the categories as ordered integers (mostly wrong)
  - Python scikit-learn: LabelEncoder
  - Ok for tree-based methods
- One Hot Encoding
  - Transform categories into individual binary (0 or 1) features
  - Python scikit-learn: DictVectorizer, OneHotEncoder
  - Ok for K-means, Linear, NNs, etc.

# Feature Encoding

- Labeled Encoding

A	0
B	1
C	2

Feature 1	Encoded Feature 1
A	0
A	0
A	0
A	0
B	1
B	1
B	1
C	2
C	2

# Feature Encoding

- One Hot Encoding

A	<b>1</b>	0	0
B	0	<b>1</b>	0
C	0	0	<b>1</b>

Feature	Feature = A	Feature = B	Feature = C
A	<b>1</b>	0	0
A	<b>1</b>	0	0
A	<b>1</b>	0	0
A	<b>1</b>	0	0
B	0	<b>1</b>	0
B	0	<b>1</b>	0
B	0	<b>1</b>	0
C	0	0	<b>1</b>
C	0	0	<b>1</b>

# Feature Encoding

- Frequency Encoding
  - Encoding of categorical levels of feature to values between 0 and 1 based on their relative frequency

A	0.44 (4 out of 9)
B	0.33 (3 out of 9)
C	0.22 (2 out of 9)

Feature	Encoded Feature
A	0.44
A	0.44
A	0.44
A	0.44
B	0.33
B	0.33
B	0.33
C	0.22
C	0.22

# Feature Encoding - Target mean encoding

- Instead of dummy encoding of categorical variables and increasing the number of features we can encode each level as the mean of the response.

A	0.75 (3 out of 4)
B	0.66 (2 out of 3)
C	1.00 (2 out of 2)

Feature	Outcome	MeanEncode
A	1	0.75
A	0	0.75
A	1	0.75
A	1	0.75
B	1	0.66
B	1	0.66
B	0	0.66
C	1	1.00
C	1	1.00

# Feature Encoding - Target mean encoding

- Also it is better to calculate weighted average of the overall mean of the training set and the mean of the level:

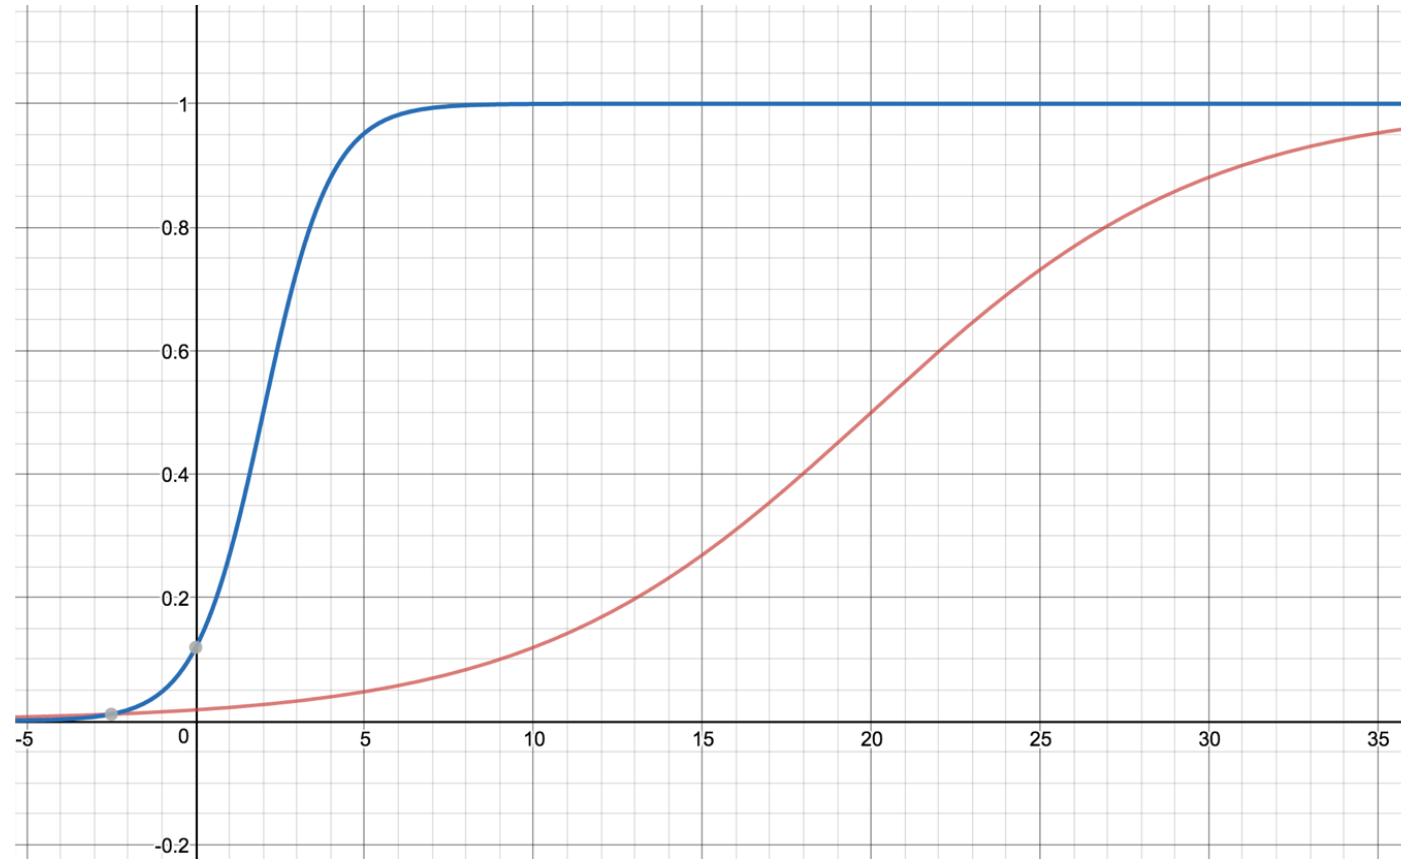
$$\lambda(n) * mean(level) + (1 - \lambda(n)) * mean(dataset)$$

- The weights are based on the frequency of the levels i.e. if a category only appears a few times in the dataset then its encoded value will be close to the overall mean instead of the mean of that level.

# Feature Encoding – Target mean encoding $\lambda(n)$ example

$$\frac{1}{1 + \exp\left(\frac{-(x-k)}{f}\right)}$$

x = frequency  
k = inflection  
point  
f = steepness



# Feature Encoding - Target mean encoding - Smoothing

$$\lambda = \frac{1}{1 + \exp(-\frac{(x - 2)}{0.25})}$$

	x	level	dataset	$\lambda$	
A	4	0.75	0.77	0.99	$0.99*0.75 + 0.01*0.77 = 0.7502$
B	3	0.66	0.77	0.98	$0.98*0.66 + 0.02*0.77 = 0.6622$
C	2	1.00	0.77	0.5	$0.5*1.0 + 0.5*0.77 = 0.885$

$$\lambda = \frac{1}{1 + \exp(-\frac{(x - 3)}{0.25})}$$

	x	level	dataset	$\lambda$	
A	4	0.75	0.77	0.98	$0.98*0.75 + 0.01*0.77 = 0.7427$
B	3	0.66	0.77	0.5	$0.5*0.66 + 0.5*0.77 = 0.715$
C	2	1.00	0.77	0.017	$0.017*1.0 + 0.983*0.77 = 0.773$

Feature	Outcome
A	1
A	0
A	1
A	1
B	1
B	1
B	0
C	1
C	1



# Feature Encoding - Target mean encoding

- Instead of dummy encoding of categorical variables and increasing the number of features we can encode each level as the mean of the response.

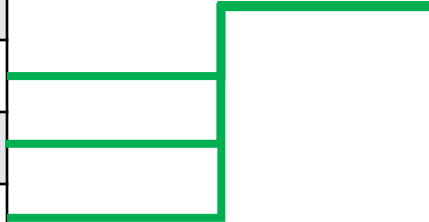
A	0.75 (3 out of 4)
B	0.66 (2 out of 3)
C	1.00 (2 out of 2)

Feature	Outcome	MeanEncode
A	1	0.75
A	0	0.75
A	1	0.75
A	1	0.75
B	1	0.66
B	1	0.66
B	0	0.66
C	1	1.00
C	1	1.00

# Feature Encoding - Target mean encoding

- To avoid overfitting we could use leave-one-out schema

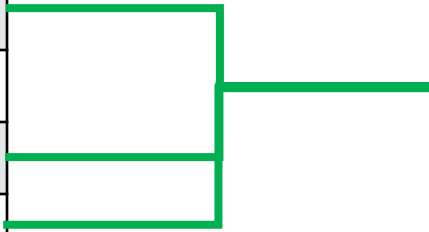
Feature	Outcome		LOOencode
A	1		0.66
A	0		
A	1		
A	1		
B	1		
B	1		
B	0		
C	1		
C	1		



The diagram illustrates the leave-one-out (LOO) encoding process for Feature A. It shows a table with columns 'Feature' and 'Outcome'. The first row has Feature A and Outcome 1. The next three rows also have Feature A, but with Outcomes 0, 1, and 1 respectively. These three rows are highlighted with green text. To the right of the table, a box labeled 'LOOencode' contains the value 0.66. Green lines connect the three rows with Outcome 1 for Feature A to the LOOencode box, indicating that the encoding is calculated as the mean of the outcomes for that feature when one instance is left out.

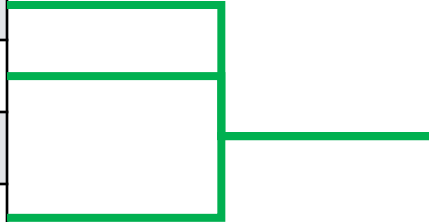
# Feature Encoding - Target mean encoding

- To avoid overfitting we could use leave-one-out schema

Feature	Outcome		LOOencode
A	1		0.66
A	0		1.00
A	1		
A	1		
B	1		
B	1		
B	0		
C	1		
C	1		

# Feature Encoding - Target mean encoding

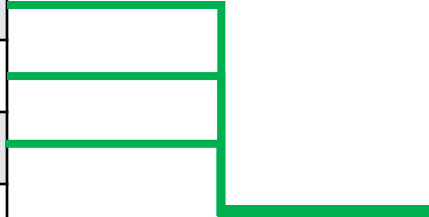
- To avoid overfitting we could use leave-one-out schema

Feature	Outcome		LOOencode
A	1		0.66
A	0		1.00
A	1		0.66
A	1		
B	1		
B	1		
B	0		
C	1		
C	1		

# Feature Encoding - Target mean encoding

- To avoid overfitting we could use leave-one-out schema


Feature	Outcome		LOOencode
A	1		0.66
A	0		1.00
A	1		0.66
A	1		0.66
B	1		
B	1		
B	0		
C	1		
C	1		



# Feature Encoding - Target mean encoding

- To avoid overfitting we could use leave-one-out schema

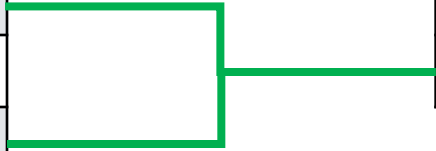
Feature	Outcome	LOOencode
A	1	0.66
A	0	1.00
A	1	0.66
A	1	0.66
B	1	0.50
<b>B</b>	<b>1</b>	
<b>B</b>	<b>0</b>	
C	1	
C	1	



# Feature Encoding - Target mean encoding

- To avoid overfitting we could use leave-one-out schema

Feature	Outcome		LOOencode
A	1		0.66
A	0		1.00
A	1		0.66
A	1		0.66
<b>B</b>	<b>1</b>		0.50
B	1		0.50
<b>B</b>	<b>0</b>		
C	1		
C	1		



# Feature Encoding - Target mean encoding

- To avoid overfitting we could use leave-one-out schema


Feature	Outcome	LOOencode
A	1	0.66
A	0	1.00
A	1	0.66
A	1	0.66
B	1	0.50
B	1	0.50
B	0	1.00
C	1	
C	1	



# Feature Encoding - Target mean encoding

- To avoid overfitting we could use leave-one-out schema


Feature	Outcome	LOOencode
A	1	0.66
A	0	1.00
A	1	0.66
A	1	0.66
B	1	0.50
B	1	0.50
B	0	1.00
C	1	1.00
C	1	



# Feature Encoding - Target mean encoding

- To avoid overfitting we could use leave-one-out schema

Feature	Outcome	LOOencode
A	1	0.66
A	0	1.00
A	1	0.66
A	1	0.66
B	1	0.50
B	1	0.50
B	0	1.00
C	1	1.00
C	1	1.00



# Feature Encoding – Weight of Evidence

$$WoE = \ln\left(\frac{\% \text{ non-events}}{\% \text{ events}}\right)$$

To avoid division by zero

$$WoE_{adj} = \ln\left(\frac{\text{Number of non-events in a group} + 0.5}{\text{Number of events in a group} + 0.5} \cdot \frac{\text{Number of events}}{\text{Number of non-events}}\right)$$

# Feature Encoding – Weight of Evidence

	Non-events	Events	% of non-events	% of events	WoE
A	1	3	50	42	$\ln\left(\frac{(1 + 0.5)/_2}{(3 + 0.5)/_7}\right) = 0.4$
B	1	2	50	29	$\ln\left(\frac{(1 + 0.5)/_2}{(2 + 0.5)/_7}\right) = 0.74$
C	0	2	0	29	$\ln\left(\frac{(0 + 0.5)/_2}{(2 + 0.5)/_7}\right) = -0.35$

Feature	Outcome	WoE
A	1	0.4
A	0	0.4
A	1	0.4
A	1	0.4
B	1	0.74
B	1	0.74
B	0	0.74
C	1	-0.35
C	1	-0.35

# Feature Encoding – Weight of Evidence and Information Value

$$IV = \sum (\% non - events - \% events) * WoE$$

	Non-events	Events	% of non-events	% of events	WoE	IV
A	1	3	50	42	$\ln\left(\frac{(1 + 0.5)/2}{(3 + 0.5)/7}\right) = 0.4$	$(0.5 - 0.42) * 0.4 = 0.032$
B	1	2	50	29	$\ln\left(\frac{(1 + 0.5)/2}{(2 + 0.5)/7}\right) = 0.74$	$(0.5 - 0.29) * 0.4 = 0.084$
C	0	2	0	29	$\ln\left(\frac{(0 + 0.5)/2}{(2 + 0.5)/7}\right) = -0.35$	$(0 - 0.29) * -0.35 = 0.105$
						0.221

Feature	Outcome	WoE
A	1	0.4
A	0	0.4
A	1	0.4
A	1	0.4
B	1	0.74
B	1	0.74
B	0	0.74
C	1	-0.35
C	1	-0.35

# Feature Encoding – Weight of Evidence and Information Value



Information Value	Variable Predictiveness
Less than 0.02	Not useful for prediction
0.02 to 0.1	Weak predictive Power
0.1 to 0.3	Medium predictive Power
0.3 to 0.5	Strong predictive Power
>0.5	Suspicious Predictive Power

# Feature Encoding – Numerical features

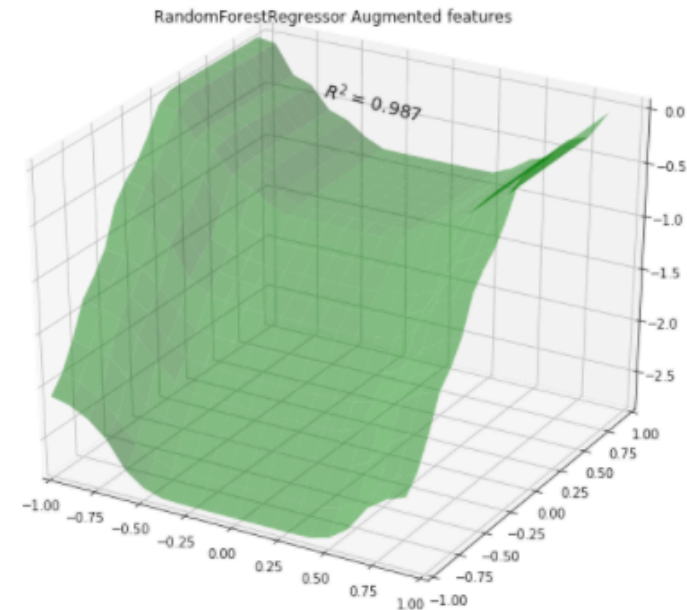
- Binning using quantiles (population of the same size in each bin) or histograms (bins of same size)
  - Replace with bin's mean or median
  - Treat bin id as a category level and use any categorical encoding schema
- Dimensionality reduction techniques – SVD and PCA
- Clustering and using cluster IDs or/and distances to cluster centers as new features

# Feature Interaction

- $y = x_1^2 - x_2^2 + x_2 - 1$



Adding  $x_1^2$  and  $x_2^2$   
as new features





## Feature Interaction – how to find?

- Domain knowledge
- ML algorithm behavior (for example analyzing GBM splits or linear regressor weights)

## Feature Interaction – how to model?

- Clustering and kNN for numerical features
- Target encoding for pairs (or even triplets and etc.) of categorical features
- Encode categorical features by stats of numerical features

# Feature Extraction

- There usually have some meaningful features inside existing features, you need to extract them manually
- Some examples
  - Location
    - Address, city, state and zip code .... (categorical or numeric)
  - Time
    - Year, month, day, hour, minute, time ranges, .... (numeric)
    - Weekdays or weekend (binary)
    - Morning, noon, afternoon, evening, ... (categorical)
  - Numbers
    - Turn age numbers into ranges (ordinal or categorical)

# Feature Extraction: Textual Data

- Bag-of-Words: extract tokens from text and use their occurrences (or TF/IDF weights) as features
- Require some NLP techniques to aggregate token counts more precisely
  - Split token into sub-tokens by delimiters or case changes
  - N-grams at word (often 2-5 grams) or character level
  - Stemming for English words
  - Remove stop words (not always necessary)
  - Convert all words to lower case
- Bag-of-Words Tools
  - Python: CountVectorizer, TfidfTransformer in scikit-learn package

# Feature Extraction: Textual Data

- Deep Learning for textual data
  - Turn each token into a vector of predefined size
  - Help compute “semantic distance” between tokens/words
    - For example, the semantic distance between user query and product titles in search results (how relevant?)
  - Greatly reduce the number of text features used for training
    - Use average vector of all words in given text
    - Vector size: 100~300 is often enough
- Tools
  - Word2vec, Doc2vec, GloVe

# End